



How to build an Al-augmented bid practice

A practical guide for business leaders developing orchestrated teams of people and specialist AI agents.

BUILD IN SEQUENCE. CONNECT PROGRESSIVELY.

1

Business outcome

2

Trusted context

3

Agents and tools

4

Evaluate and refine

5

Human decision gates

6

Measure and learn





Why now is the time to begin

AI agents are moving the conversation beyond faster drafting. They can hold a goal, use tools, retrieve context and complete bounded steps across a workflow. Orchestration adds the missing management layer: it breaks a larger outcome into work, routes that work to specialist agents and brings it back for evaluation and human decision.

That can sound like a major transformation programme. It does not have to begin that way. The most credible route is a controlled pilot in one real workflow: useful enough to matter, bounded enough to understand and safe enough to stop. The first goal is not autonomy. It is dependable assistance inside a process the business can see and govern.

Each pilot should leave reusable assets behind: a clearer process, cleaner sources, explicit permissions, a test set, decision gates and evidence about what changed. Add the next specialist only when the first one earns trust. Connect the workflow as the foundations mature.

THE CENTRAL IDEA

**Start with one
useful agent.
Earn the right
to connect
the next.**

What this guide covers

- | | |
|---|--|
| 01 Business outcome
Choose the problem before the technology. | 02 Trusted context
Prepare the evidence, access and rules. |
| 03 Agents and tools
Design a small team of clear specialist roles. | 04 Evaluate and refine
Test outputs, exceptions and failure modes. |
| 05 Human decision gates
Make authority, approval and recovery explicit. | 06 Measure and learn
Expand only from evidence, not enthusiasm. |



Six pillars, built in sequence

Think of the system as a managed practice, not one clever assistant. The business sets the outcome. Trusted context forms the foundation. Specialist agents and tools do bounded work. Evaluation, human authority and learning keep the whole system dependable.



Direction of build and evidence

The building analogy

A strong AI-augmented practice is built like a house. The business outcome is the brief. Trusted knowledge, permissions and integrations are the foundations and services. Specialist agents are the trades. Evaluation is inspection. Human decision gates are sign-off. Measurement becomes the snagging list and the blueprint for the next build.

The analogy matters because attractive automation cannot compensate for weak foundations. If sources conflict, access is unreliable or authority is vague, adding more agents usually multiplies the uncertainty.

What orchestration does

- Decomposes one business outcome into clear work packages.
- Selects the right specialist, system or person for each step.
- Maintains state, evidence and dependencies across handovers.
- Surfaces uncertainty and exceptions instead of hiding them.

A useful orchestrator resembles an air-traffic controller: it coordinates specialists, keeps the whole system visible and intervenes when the planned route no longer fits reality.



Begin with the job and its foundations

1 Business outcome

Define the decision or result the system must improve.

Do not begin by asking where an AI agent could be used. Begin with a recurring business constraint: an opportunity spotted too late, a qualification decision made with incomplete evidence, a slow mobilisation or avoidable rework before submission. Describe the outcome in operational terms, then choose the smallest part of the workflow that could improve it.

WHAT TO PUT IN PLACE

- Map the present workflow, including waits, handovers and rework.
- Name the consequential decision, its owner and its acceptance criteria.
- Set a baseline for quality, cycle time, effort, risk or capacity.
- Choose a bounded pilot where success and failure will both be visible.

LEADER'S QUESTIONS

What must improve? • Who owns the decision? • What would make us stop?

2 Trusted context

Give the system a clean, controlled working environment.

Agents need more than access to documents. They need an agreed source hierarchy, current information, permissions and a way to show where an answer came from. Treat context as a deliberately prepared workspace, not a data lake. Build a small reliable foundation for the pilot before connecting the wider estate.

WHAT TO PUT IN PLACE

- Define approved sources and the order of precedence when they conflict.
- Remove duplicates, expired evidence and ambiguous ownership.
- Give each role only the access and tools needed for its task.
- Record source, version, retrieval date and evidence trail with every output.

LEADER'S QUESTIONS

What can it trust? • What may it access? • How will we trace the answer?



Build a small team and test it hard

3 Agents and tools

Give every digital specialist a bounded role.

A specialist agent should have a clear job, known inputs, permitted tools, a defined output and an escalation route. Think of a building team: an electrician is valuable because the role is specific, not because one trade tries to do everything. Start with one specialist. Add an orchestrator only when real coordination between roles creates value.

WHAT TO PUT IN PLACE

- Write a role card: purpose, inputs, tools, output, limits and escalation.
- Use APIs for structured actions and controlled access to systems.
- Keep memory purposeful: enough state to continue work, never an uncontrolled archive.
- Prefer deterministic automation where a fixed rule is safer and simpler.

LEADER'S QUESTIONS

Why an agent? • What can it do? • When must it hand back?

4 Evaluate and refine

Make performance visible before work advances.

A plausible answer is not the same as a dependable output. Build evaluation into the workflow. Test the sources used, accuracy, completeness, format, uncertainty and known exceptions. Run the pilot in shadow mode against representative historical and live cases. Record failures without smoothing them away, then refine the role, context or control that caused them.

WHAT TO PUT IN PLACE

- Create a test set that includes ordinary work, edge cases and bad inputs.
- Compare the agent with the current baseline, not with a perfect ideal.
- Use independent checks for evidence, calculation, compliance and completeness.
- Separate model error from weak process, missing context and unclear instructions.

LEADER'S QUESTIONS

What is tested? • Which failures matter? • Is review effort actually lower?



Keep people in command

5 Human decision gates

Match delegated authority to consequence and reversibility.

Human control should be designed, not added as a final approval button. Low-risk, reversible steps may run automatically inside policy. Commercial, contractual, compliance and submission decisions need explicit ownership and evidence. Like an aircraft autopilot, the system can handle routine work while people retain situational awareness, intervene and take control when conditions change.

WHAT TO PUT IN PLACE

- Classify actions by consequence, reversibility and confidence required.
- Set approval thresholds, decision owners and service times for each gate.
- Show what the system did, which evidence it used and where it is uncertain.
- Provide stop, correct, escalate and recover routes before live operation.

LEADER'S QUESTIONS

Who decides? • What is reversible? • How do we recover safely?

6 Measure and learn

Promote the workflow only when the evidence supports it.

Measure the whole operating result, not just the speed of one generated task. Track decision quality, rework, review burden, cycle time, capacity released, exceptions and downstream outcomes. Use corrections as governed learning inputs. Do not allow an agent to absorb every human edit automatically: first decide whether the edit fixes an error, reflects preference or introduces a new risk.

WHAT TO PUT IN PLACE

- Review leading indicators during the pilot and outcome indicators over time.
- Compare quality and control alongside time or cost saved.
- Turn repeated corrections into updates to source, role, test or policy.
- Choose deliberately: stop, refine, repeat, connect another role or scale.

LEADER'S QUESTIONS

What changed end to end? • What did review cost? • Has the system earned more scope?



Design one controlled pilot

The best first pilot is neither trivial nor transformational. It is a Goldilocks use case: repeatable enough to structure, important enough to matter and bounded enough to evaluate. A useful pattern is a high-volume workflow that currently depends on people moving information between documents and systems.



CHOOSE

One workflow

Name the bottleneck, outcome, owner and baseline. Exclude any decision whose consequence cannot yet be governed.



PREPARE

A trusted workspace

Assemble approved sources, permissions, rules, expected outputs and representative cases.



BUILD

One specialist

Give the agent one bounded role and the minimum tools it needs. Keep the first handover to a person.



TEST

Shadow mode

Run beside the existing process. Capture accuracy, completeness, exceptions, review effort and failure recovery.



CONNECT

A second role

Add a second specialist only when the handover between roles creates measurable value.



ORCHESTRATE

The workflow

Coordinate roles when sequencing, shared state and routing become the constraint.

A pilot succeeds when leaders can make a better next decision: stop, refine, repeat, connect or scale.



Start small. Build deliberately.

The goal is not a collection of AI tools. It is a working bid practice in which people and digital specialists share a clear outcome, trusted context and explicit decision rights. The route there is progressive: prove one role, strengthen the foundations and connect the next part of the workflow only when the evidence justifies it.

A practical first 90 days

WEEKS 1-2**Frame**

Choose one workflow. Map the current work, business outcome, owner, baseline and non-negotiable controls.

WEEKS 3-4**Prepare**

Create the trusted workspace, role card, test cases, approval thresholds and recovery path.

WEEKS 5-8**Build and test**

Run one specialist agent in shadow mode. Examine failures and review effort, not just impressive outputs.

WEEKS 9-12**Decide**

Review the evidence with the business sponsor and people doing the work. Stop, refine, repeat or connect the next role.

READY TO MAP YOUR FIRST WORKFLOW?

Begin with the business decision, then design the smallest credible pilot around it.

ignisleadership.com



Build or transform the practice behind complex bids.

Better bid practices, designed for people and AI agents.

Bid Management Consultancy & Advisory

ignisleadership.com